

# OpenCL introduction

# OpenCL

- Data and task parallel model
- Derived from the ISO C99 standard
  - With parallel extensions
- Numeric operations performed according to the IEEE754 standard
- Support of embedded and mobile devices
- Data transfer between OpenGL, OpenGL ES

See OpenCL specification for details

# OpenCL


- Heterogeneous platform support
  - Parallel CPU cores
  - GPU
  - Digital Signal Processor (DSP)
  - Cell/B.E. processor

# OpenCL Architecture

- Elements of the OpenCL architecture
  - Platform model
  - Execution model
  - Memory model
  - Programming model

# Platform model

- Host device
- OpenCL device
- Computing Unit (CU)
  - Processing Element (PE)
    - Single Instruction Multiple Data SIMD (common program counter)
    - Single Program Multiple Data SPMD (independent program counters)



# Execution model

- Host program
  - Context management
  - Execution control
- Kernel program
  - Controlling the CUs

# Execution model

- Kernel program
  - Index space (NDRange)
  - Work-groups
  - Work-items
    - global ID
    - Same programs in the work-group
    - The execution control can differ in different units

# Execution model

- Kernel program
  - Index space (NDRange)
  - Work-groups
    - Finer indexing mechanism
    - Work-group ID
    - Local ID for the Work-Items
  - Work-Items



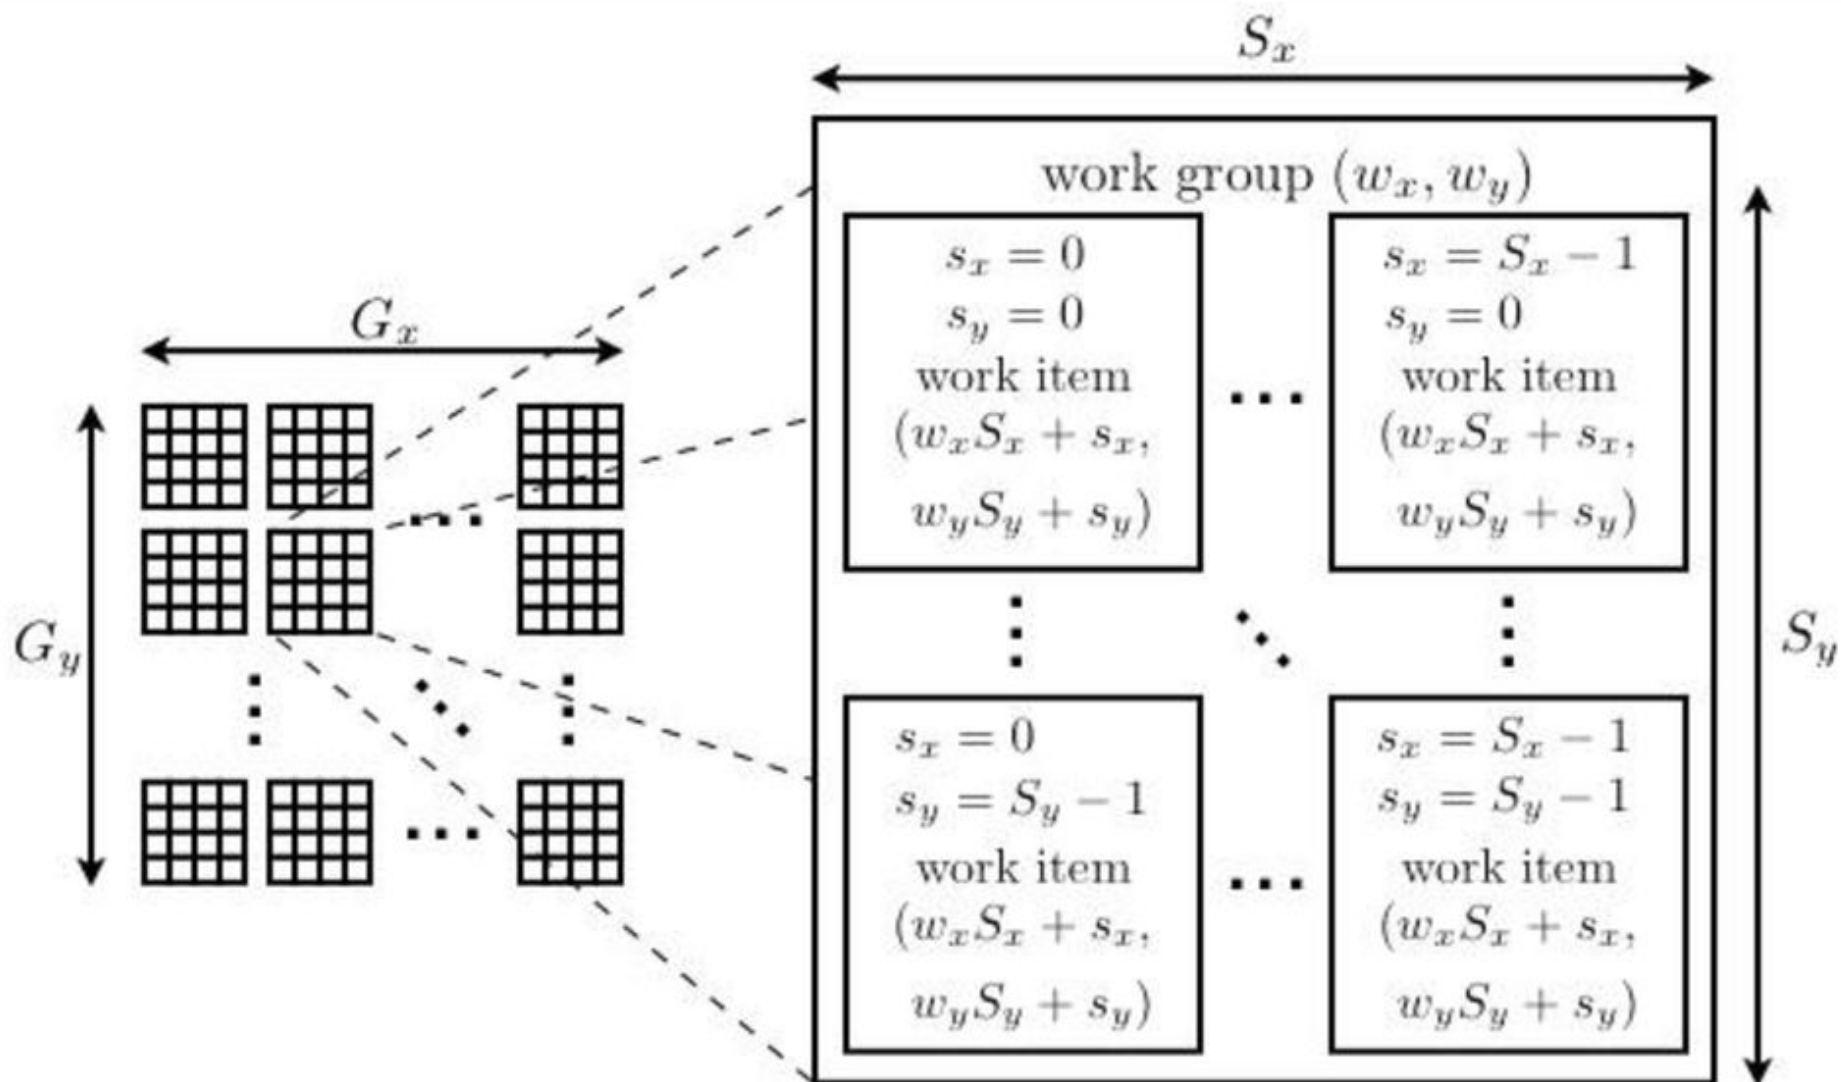
# Execution model

- Kernel program
  - Index space (NDRange)
    - N dimensional problem space ( $N=1,2,3$ )
    - Each index has the same dimensionality
  - Indexing
    - Global index space:  $(G_x, G_y)$
    - Size of work-groups:  $(S_x, S_y)$ 
      - Work-group ID  $(w_x, w_y)$
    - Local ID  $(s_x, s_y)$

# Execution model

- Calculating IDs
  - Global address space:  $(G_x, G_y)$
  - Work-group size:  $(S_x, S_y)$ 
    - Work-group ID:  $(w_x, w_y)$
  - Local ID:  $(s_x, s_y)$
- Global ID:  $(g_x, g_y) = (w_x \cdot S_x + s_x, w_y \cdot S_y + s_y)$
- Number of work-groups:  $(W_x, W_y) = (G_x / S_x, G_y / S_y)$
- Work-group ID:  $(w_x, w_y) = ((g_x - s_x) / S_x, (g_y - s_y) / S_y)$

# Execution model



# Execution model

## ■ Context

- Devices: a set of OpenCL capable devices
- Kernels: a set of OpenCL functions
- Program objects:
  - Kernel source code
  - Executable binary representation
- Memory objects:
  - Memory used by the host and the OpenCL devices
  - Other values seen by the kernels

# Execution model

- Command-queue
  - A command stream controlled by the host
  - Controls the execution of the threads
  - Commands:
    - Kernel execution
    - Memory operations
    - Synchronization

# Execution model

- Command-queue execution modes
  - In-order execution
    - FIFO
    - Serializes the execution order of commands in a queue
    - a prior command on the queue completes before the following command begins
  - Out-of-order execution
    - Commands are issued in order, but do not wait to complete before following commands execute
    - Any order constraints are enforced by the programmer through explicit synchronization commands.



# Execution model

- Kernel types
  - OpenCL kernel
    - OpenCL C functions
    - Executable on an OpenCL device
  - Native kernel
    - Functions accessed through host function pointer
    - Can share memory objects with OpenCL kernels
    - Optional support

# Memory model

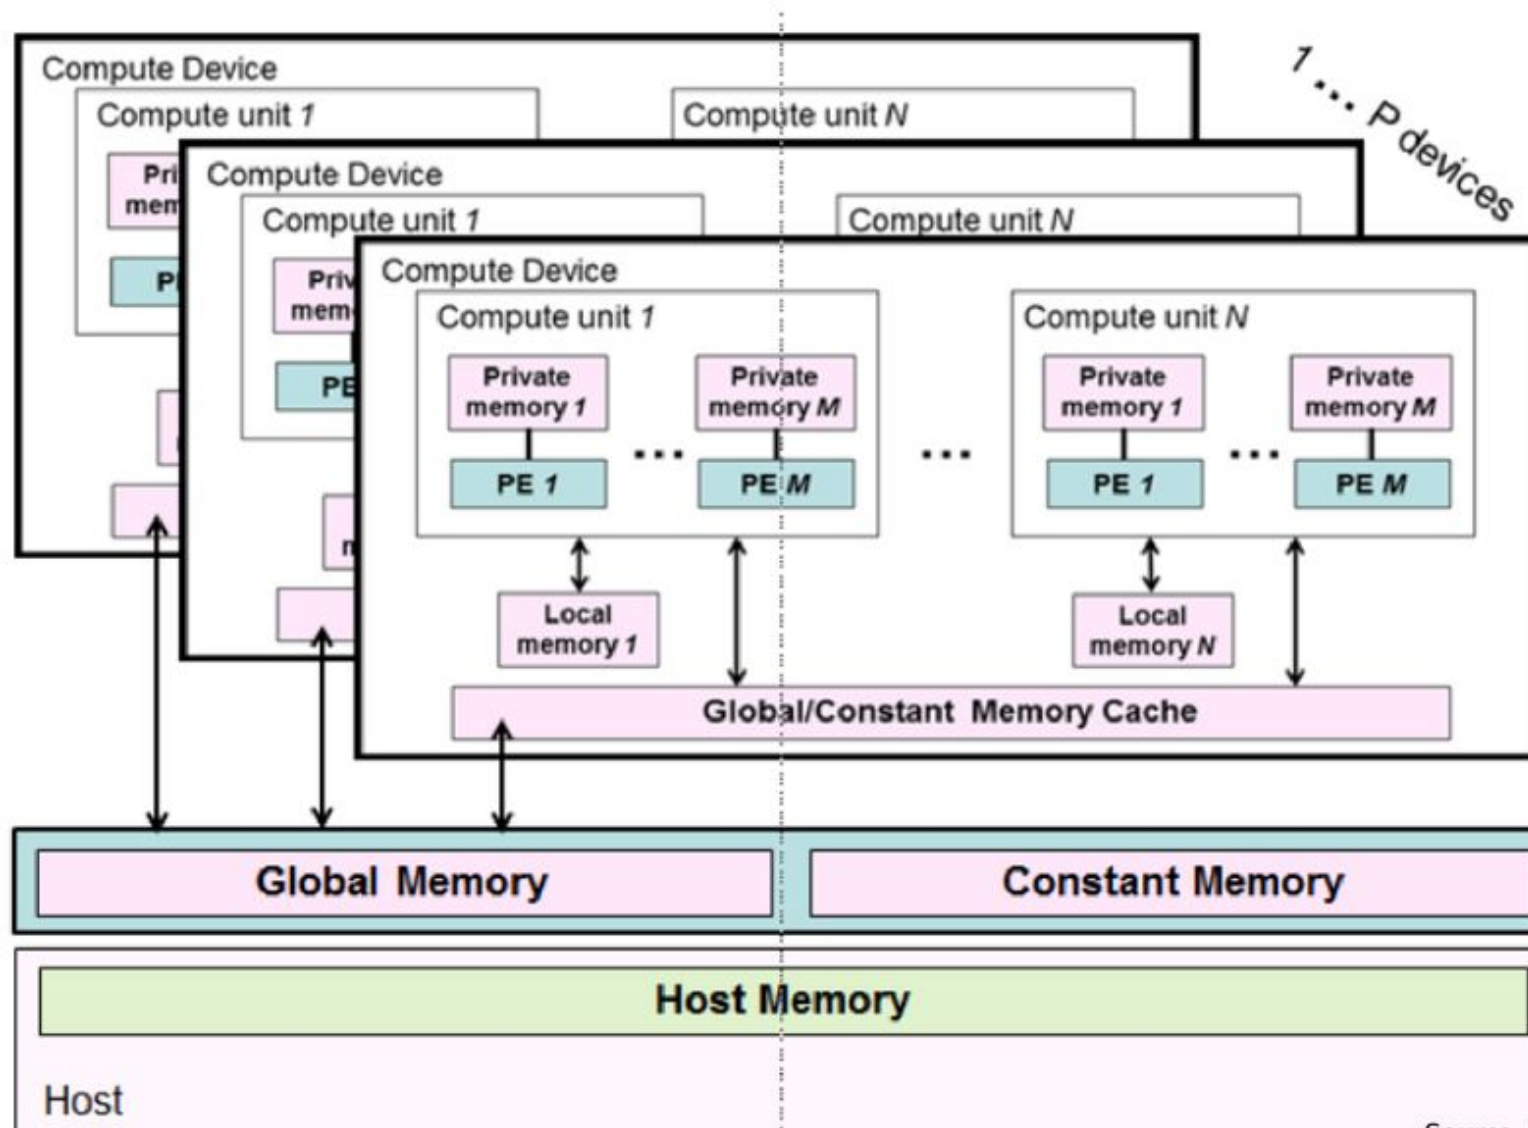
- Four distinct memory regions are available
  - Global memory
    - Work-items can read or write any element of it
    - Allocated by the host
  - Constant memory
    - Remains constant during the execution of the kernel
    - The host allocates and initializes it
    - Can be defined statically in the kernel



# Memory model

- Four distinct memory regions are available
  - Local memory
    - Shared memory within a work-group
    - Each work-item in a work-group can read or write it
    - Not visible from the host
    - May be implemented as dedicated regions of memory
  - Private memory
    - A region of memory private to a work-item
    - Seen only by the work-item

# Memory model



# Memory model

- The global memory is handled by the host
  - Memory allocation
  - Copy data to the memory objects
    - Synchronous and asynchronous operations
  - Release of memory objects
- The global memory can be mapped into the host memory address space

# Memory model

- Relaxed Consistency
  - The state of memory, visible to a work-item, is not guaranteed to be consistent across the collection of work-items at all times
  - Within a work-item memory has load / store consistency
  - Local memory
    - Consistent in a single work-group
  - Global memory
    - Consistent in a single work-group
    - No guarantees of memory consistency between different work-groups
  - Consistency for memory objects shared between enqueued commands is enforced at a synchronization point

# Programming model

- Data parallel model
  - Defines a computation in terms of a sequence of instructions applied to multiple elements of a memory object
  - The index space defines the work-items and how the data maps on the work-items
  - Not restricted to one-to-one mapping
- Hierarchical data parallelism
  - Explicit model
    - Total number of work-items and their division into work-groups
  - Implicit model
    - Only the number of work-items is specified the division into work-groups is automatic



# Programming model

- Task parallel model
  - The kernel has only one instance
  - Independent from the index space
  - Operations on vector types
  - Multiple independent tasks
  - Enqueueing native kernels to run them in parallel

# Synchronization

- Synchronization within a work-group
  - Synchronize work-items
  - work-group barrier
    - Blocking call
    - Each work-item has to reach the barrier before any are allowed to continue
  - There is no mechanism for synchronization between work-groups

# Synchronization

- Synchronization in a command-queue
  - In case of out-of-order execution
  - command-queue barrier
    - Ensures that all previously queued commands have finished execution
    - Resulting updates to memory objects are visible to subsequently enqueued commands
    - Cannot be used to synchronize between command-queues
  - Waiting on an event
    - Each function generates an event that identifies the command and memory objects it updates
    - The execution of command can be suspended until the occurrence of some events



# OpenCL C

- Scalar types
  - bool
  - unsigned char, char (8 bit integer)
  - unsigned short, short (16 bit integer)
  - unsigned int, int (32 bit integer)
  - unsigned long, long (64 bit integer)
  - float (IEEE754 floating-point)
  - half (16 bit float)
  - size\_t (return type of the sizeof operator 32/64 bit)
  - ptrdiff\_t (difference between two pointers 32/64 bit)
  - (u)intptr\_t (pointer type)
  - void

# OpenCL C

- Vector types

- (u)char*n*
- (u)short*n*
- (u)int*n*
- (u)long*n*
- float*n*

- The signed values represented in two's complement form
- (u) stands for unsigned
- *n* can be 2,4,8,16

# OpenCL C

- Vector components
  - Swizzle operator (.xyzw)
    - `float4 f; f.xy; f.xxyy;`
  - Numeric indices (.s[o-g|a-f|A-F])
    - `float4 f; f.s12;`
    - `float16; f.saBcdE`
  - Halving (.odd, .even, .lo, .hi)
    - `float4 f; f.hi; f.even.lo;`
    - `float4 left, right;`  
`float8 interleaved;`  
`interleaved.even = left; interleaved.odd = right;`

# OpenCL C

- Conversion between different types
  - Implicit conversion
    - Limited usability
    - Between scalar types
  - Explicit conversion
    - Scalar – Vector conversion
      - `float4 f = (float4)1.0;`
    - Conversion between vector types
      - `destType convert_destType_sat_roundingMode(sourceType)`
        - `_sat` – truncation to the codomain
        - `_roundingMode` – rounding
      - `uchar4 u; int4 c = convert_int4(u);`

# OpenCL C

- Conversion between types
  - Types should have same size
  - `as_type()`
    - `float f = 1.0f;`  
`uint u = as_uint(f); // the value will be: 0x3f800000`
    - `float4 f = (float4)(1.0f, 2.0f, 3.0f, 4.0f);`  
`int4 i = as_int4(f);`  
`// (0x3f800000, 0x40000000, 0x40400000, 0x40800000)`

# OpenCL C

- Address space qualifiers
  - `__global` : global memory
    - `__global float4 color;`
  - `__local` : local memory
    - `__local float16 shared;`
  - `__constant` : constant memory
    - `__constant float uniformData;`
    - Can be initialized from the host
  - `__private` : private memory
    - `__private float8 workItemExclusive;`



# OpenCL C

- Function qualifiers
  - `__kernel` : OpenCL function
    - Only an OpenCL device can execute it
    - The host program can call it
    - Other OpenCL kernels can call it
  - `__attribute__` : hints to the compiler
    - `vec_type_hint(typen)` : size of vector operations
      - Work-items can be merged or separated by the compiler to better match the hardware capabilities

# OpenCL C

- Built-in functions

- Work-item information:

- `uint get_work_dim()`
    - `size_t get_global_size(uint dimIdx);`
    - `size_t get_global_id(uint dimIdx);`
    - `size_t get_local_size(uint dimIdx);`
    - `size_t get_local_id(uint dimIdx);`
    - `size_t get_num_groups(uint dimIdx);`
    - `size_t get_group_id(uint dimIdx);`



# OpenCL C

- Built-in functions
  - Math functions
    - E.g. sin, cos, tan, floor ...
    - float, half, integer types
  - Common functions
    - E.g. clamp, min, max ...
    - float types
  - Geometric functions
    - E.g. cross, dot, length, normalize ...
    - float types
  - Relational functions
    - E.g. isequal(floatn, floatn)  
isfinite(float)
    - float types

# OpenCL C

- Built-in functions
  - Vector load functions
    - pointer – vector conversion
  - Vector store functions
    - Vector – pointer conversion

# OpenCL C

- Built-in functions
  - Synchronization functions
    - `barrier(flag);`
      - All work-items in a work-group must execute this function before any are allowed to continue
      - `CLK_LOCAL_MEM_FENCE` : makes the local memory consistent
      - `CLK_GLOBAL_MEM_FENCE` : makes the global memory consistent
    - `mem_fence(flag);`
      - Loads and stores will be committed to memory
    - `read_mem_fence(flag);`
    - `write_mem_fence(flag);`

# OpenCL C

- Built-in functions
  - Async Copy functions
    - From global memory to local memory
    - From local memory to global memory
    - `event_t async_work_group_copy(...);`
    - `wait_group_events(..., eventList);`
  - Prefetch
    - Loads a part of the global memory to the cache
  - Atomic functions
    - E.g. `atomic_add`

# My first OpenCL program

```
#include <iostream>
#include <CL/opencl.h>

#define DATA_SIZE (1024*1240)

int main(int argc, char* argv[]){
    cl_int err;

    size_t global; // global space
    size_t local;  // local space

    cl_platform_id platform;
    err = clGetPlatformIDs(1, &platform, NULL);
    if(err != CL_SUCCESS){
        std::cerr << "Error: Failed to find a platform!" << std::endl;
        return EXIT_FAILURE;
    }

    //...
```



# My first OpenCL program

```
cl_device_id device_id;
err = clGetDeviceIDs(platform, CL_DEVICE_TYPE_GPU, 1, &device_id, NULL);
if(err != CL_SUCCESS){
    std::cerr << "Error: Failed to create a device group!" << std::endl;
    return EXIT_FAILURE;
}

cl_context context;
context = clCreateContext(0, 1, &device_id, NULL, NULL, &err);
if (!context) {
    std::cerr << "Error: Failed to create a compute context!" << std::endl;
    return EXIT_FAILURE;
}

cl_command_queue commands;
commands = clCreateCommandQueue(context, device_id, 0, &err);
if (!commands) {
    std::cerr << "Error: Failed to create a command commands!" << std::endl;
    return EXIT_FAILURE;
}
// ...
```

# My first OpenCL program

```
cl_program program;
program = clCreateProgramWithSource(context, 1,
                                   (const char **) &KernelSource, NULL, &err);

if (!program) {
    std::cerr << "Error: Failed to create compute program!" << std::endl;
    return EXIT_FAILURE;
}

err = clBuildProgram(program, 0, NULL, NULL, NULL, NULL);
if (err != CL_SUCCESS) {
    size_t len;
    char buffer[2048];

    std::cerr << "Error: Failed to build program executable!" << std::endl;
    clGetProgramBuildInfo(program, device_id, CL_PROGRAM_BUILD_LOG,
                          sizeof(buffer), buffer, &len);
    std::cerr << buffer << std::endl;
    exit(1);
}

// ...
```

# My first OpenCL program

```
cl_kernel kernel;
kernel = clCreateKernel(program, "square", &err);
if (!kernel || err != CL_SUCCESS) {
    std::cerr << "Error: Failed to create compute kernel!" << std::endl;
    exit(1);
}

float* data = new float[DATA_SIZE];    // input array
float* results = new float[DATA_SIZE]; // output array
unsigned int correct;
cl_mem input;                          // device memory object for the input
cl_mem output;                         // device memory object for the output

// the input values are random
unsigned int count = DATA_SIZE;
for(int i = 0; i < count; i++){
    data[i] = rand() / (float)RAND_MAX;
}

// ...
```



# My first OpenCL program

```
input = clCreateBuffer(context, CL_MEM_READ_ONLY, sizeof(float) * count,
                          NULL, NULL);
output = clCreateBuffer(context, CL_MEM_WRITE_ONLY, sizeof(float) * count,
                          NULL, NULL);

if (!input || !output) {
    std::cerr << "Error: Failed to allocate device memory!" << std::endl;
    exit(1);
}

// copy input values to the global memory of the device
err = clEnqueueWriteBuffer(commands, input,
                           CL_TRUE, 0, sizeof(float) * count,
                           data, 0, NULL, NULL);

if (err != CL_SUCCESS) {
    std::cerr << "Error: Failed to write to source array!" << std::endl;
    exit(1);
}

// ...
```

# My first OpenCL program

```
// Kernel arguments
err = 0;
err = clSetKernelArg(kernel, 0, sizeof(cl_mem), &input);
err |= clSetKernelArg(kernel, 1, sizeof(cl_mem), &output);
err |= clSetKernelArg(kernel, 2, sizeof(unsigned int), &count);
if (err != CL_SUCCESS) {
    std::cerr << "Error: Failed to set kernel arguments! " << err << std::endl;
    exit(1);
}

// Setting up the work-group size
err = clGetKernelWorkGroupInfo(kernel, device_id,
                                CL_KERNEL_WORK_GROUP_SIZE,
                                sizeof(local), &local, NULL);

if (err != CL_SUCCESS) {
    std::cerr << "Error: Failed to retrieve kernel work group info! "
                << err << std::endl;
    exit(1);
}
// ...
```

# My first OpenCL program

```
// Enqueuing the kernel
global = count;
err = clEnqueueNDRangeKernel(commands, kernel,
                             1, NULL, &global, &local,
                             0, NULL, NULL);

if (err) {
    std::cerr << "Error: Failed to execute kernel!" << std::endl;
    return EXIT_FAILURE;
}

// Waiting for the kernel to be executed
clFinish(commands);
// Reading the result
err = clEnqueueReadBuffer( commands, output,
                           CL_TRUE, 0, sizeof(float) * count,
                           results, 0, NULL, NULL );

if (err != CL_SUCCESS) {
    std::cerr << "Error: Failed to read output array! " << err << std::endl;
    exit(1);
}
// ...
```



# My first OpenCL program

```
// Validating the result
correct = 0;
for(int i = 0; i < count; i++) {
    if(results[i] == data[i] * data[i])
        correct++;
}

std::cout << "Computed " << correct << "/" <<
count << " correct values" << std::endl;
std::cout << "Computed " << 100.f * (float)correct/(float)count <<
"% correct values" << std::endl;

// ...
```

# My first OpenCL program

```
// Cleaning
delete [] data; delete [] results;

clReleaseMemObject(input);
clReleaseMemObject(output);
clReleaseProgram(program);
clReleaseKernel(kernel);
clReleaseCommandQueue(commands);
clReleaseContext(context);

return 0;

}
```

# My first OpenCL program

## ■ OpenCL kernel

```
__kernel void square(  
    __global float* input,  
    __global float* output,  
    const unsigned int count){  
    int i = get_global_id(0);  
    if(i < count){  
        output[i] = input[i] * input[i];  
    }  
}
```